

## Assignment 3

Due Monday, March 26 at 11:59pm

**General Instructions**

- Feel free to talk to other members of the class in doing the homework. You should, however, write down your solutions yourself. *List the names of everyone you worked with at the top of your submission.*
- Keep your solutions brief and clear.
- Please use Piazza if you have questions about the homework but do not post answers. Feel free to use private posts or come to the office hours.

**Homework Submission**

- We DO NOT accept late homework submissions.
- We will be using Compass for collecting the homework assignments. Please submit your answers via [Compass](#). Hard copies are not accepted.
- Contact the TAs if you are having technical difficulties in submitting the assignment; attempt to submit well in advance of the due date/time.
- The homework must be submitted in **pdf** format. Scanned handwritten and/or hand-drawn pictures in your documents won't be accepted.
- Please do not zip the answer document (PDF) so that the graders can read it directly on Compass. You need to submit one answer document, named as **hw3\_netid.pdf**.
- Please see the [assignments](#) page for more details. In particular, we will be announcing errata, if any, on this page.
- **For this HW, please use syntax and ONLY syntax that was covered in class.**

# 1 Single-Relation Queries (30 pts)

1. [10] Consider the following relation:

`Graph(n1, n2)`

A tuple (n1, n2) in Graph stores a directed edge from a node n1 to a node n2 in the corresponding graph. Your goal is to, for *every* node in the graph, count the number of outgoing edges of that node. Note that for nodes without any outgoing edges, their edge count would be zero; you need to output this as well.

You can assume that (1) there are no duplicates or null values in the table; and (2) every node in the graph is involved in at least one edge.

**ANSWER:**

```
SELECT n1 as n, COUNT(*) as cnt FROM Graph GROUP BY n1
UNION
SELECT n2 as n, 0 as cnt FROM Graph
WHERE n2 NOT IN (SELECT n1 FROM Graph)
GROUP BY n2
```

2. [10] Consider the following relation:

`Trained(student, master, year)`

A tuple (S, M, Y) in Trained specifies that a SQL Master M trained student S who graduated in year Y. Your goal is to find *the count of* distinct SQL Masters who trained a student who graduated in the same year that 'Alice' or 'Bob' graduated.

**ANSWER:**

```
SELECT COUNT(DISTINCT(master))
FROM Trained
WHERE year IN
      (SELECT year FROM Trained
       WHERE student = 'Alice' OR student = 'Bob')
```

3. [10] Consider the following relation:

`DBMS(operator, system, performance)`

A tuple (O, S, P) in DBMS specifies an operator O in system S and has the performance value P. Your goal is to find those systems whose operators achieves a higher performance value on average than the average performance value in a system named 'PostgreSQL'.

**ANSWER:**

```
SELECT system
FROM DBMS
GROUP BY system
```

```

HAVING AVG(performance)>
    (SELECT AVG(performance)
     FROM DBMS
     WHERE system = 'PostgreSQL')

```

**ALTERNATIVE ANSWER:**

```

SELECT system
FROM DBMS
GROUP BY system
HAVING AVG(performance)>
    (SELECT AVG(performance)
     FROM DBMS
     GROUP BY system
     HAVING system = 'PostgreSQL')

```

## 2 Multi-Relation Queries (20 pts)

Consider the following relations representing student information at UIUC:

```

Mentorship(mentee_sid, mentor_sid)
Study(sid, credits)
Enrollment(did, sid)
Student(sid, street, city)

```

- A tuple (M1, M2) in Mentorship specifies that M2 is a mentor of another student M1.
- A tuple (S, C) in Study specifies that the student S has taken C credits.
- A tuple in Enrollment (D, S) specifies that student S is enrolled in department D.
- A (ST, S, C) in Student specifies that student ST lives on street S in city C.

1. [10] Find all students who live in the same city and on the same street as their mentor.

**ANSWER:**

```

SELECT s1.sid
FROM Student s1, Student s2, Mentorship m
WHERE s1.sid = m.mentee_sid AND s2.sid = m.mentor_sid AND
      s1.street = s2.street AND s1.city = s2.city

```

2. [10] Find all students (sid) who have taken more credits than the average credits of all of the students of their department.

**ANSWER:**

```

SELECT DISTINCT(s1.sid)
FROM Study s1, (
    SELECT AVG(s.credits) AS avg_credits, e.did
    FROM Study s, Enrollment e
    WHERE e.sid = s.sid
    GROUP BY e.did) as s2,
    Enrollment e
WHERE s1.credits > s2. avg_credits
AND s1.sid = e.sid
AND e.did = s2.did

```

### 3 Database Manipulation and Views (25 pts)

1. [5] In the **Study** relation, insert a new student, whose id is 66666 and has 0 credits.

```

INSERT INTO Study (sid, credits)
VALUES (66666, 0);

```

2. [5] In the **Study** relation, delete students who have graduated (i.e., the ones who have more than 200 credits).

**ANSWER:**

```

DELETE FROM Study
WHERE credits > 200;

```

3. [5] In the **Study** relation, add 2 credits for students who are mentors.

**ANSWER:**

```

UPDATE Study
SET credits = credits + 2
WHERE sid IN (SELECT mentor_sid FROM Mentorship);

```

4. [10] Incoming students are those who have been accepted (i.e., exist in the **Student** relation) but have not registered in any department (i.e., do not exist in the **Enrollment** relation). Create a View that contains **sid** of all incoming students.

**ANSWER:**

```

CREATE VIEW IncomingStudents AS
    SELECT sid FROM Student
    EXCEPT
    SELECT sid FROM Enrollment

```

## 4 Constraints and Triggers (25 pts)

1. [10] Consider the following relation:

`Payment(salary, bonus)`

Write a schema-level assertion using the “CREATE ASSERTION” statement to ensure that no bonus is larger than the maximum salary in the `Payment` relation.

**ANSWER:**

```
CREATE ASSERTION BonusControl CHECK (
    NOT EXISTS (SELECT * FROM Payment WHERE bonus >
                (SELECT MAX(salary) FROM Payment)))
```

2. [15] Consider the following relation:

`Study(sid, major, GPA)`

Write a trigger T1 that increases the GPA by 10% for those students who transform their major from any Non-CS major to ‘CS’.

**ANSWER:**

```
CREATE TRIGGER T1
    AFTER UPDATE OF major ON Study
    REFERENCING
        OLD ROW AS O, NEW ROW AS N
    FOR EACH ROW
    WHEN (N.major = 'CS' AND O.major <> 'CS')
    UPDATE Study
    SET GPA = 1.1 * GPA
    WHERE sid = N.sid
```